

Machine Learning - 100 Questions

Pages 12+ Part 1: Encoding, Features, Evaluation

Nova IMS
Generated: January 15, 2026

MACHINE LEARNING - 100 MULTIPLE CHOICE QUESTIONS

Based on Pages 12+ of Curriculum Analysis

(LU2: Encoding/Feature Engineering + LU3: Evaluation + LU4: Algorithms + LU5: Projects)

LEVEL 1: CATEGORICAL ENCODING BASICS (Questions 1-15)

Q1. Why do machine learning algorithms require numerical inputs?

- A) For visual appeal
- B) Algorithms perform mathematical operations requiring numbers
- C) To save memory
- D) No specific reason

ANSWER: B

Q2. What is label encoding?

- A) Creating binary columns
- B) Assigning integers to each category (0, 1, 2, ...)
- C) Using category names as-is
- D) Removing categories

ANSWER: B

Q3. What problem does label encoding create for nominal variables?

- A) Too slow
- B) Implies ordinal relationship when none exists
- C) Uses too much memory
- D) Always accurate

ANSWER: B

Q4. When is label encoding appropriate?

- A) Always for all categorical variables
- B) For ordinal variables or tree-based models
- C) Never in any situation
- D) Only for binary variables

ANSWER: B

Q5. What is one-hot encoding?

- A) Assigning integers to categories
- B) Creating binary column for each category
- C) Removing categories
- D) Using category frequencies

ANSWER: B

Q6. What is the main advantage of one-hot encoding?

- A) Creates fewer features
- B) No ordinal assumption for nominal data
- C) Faster computation
- D) Works only with trees

ANSWER: B

Q7. What is the main disadvantage of one-hot encoding?

- A) Too slow to compute
- B) Increases dimensionality (creates many columns)
- C) Only works for binary variables

D) Requires target variable

ANSWER: B

Q8. For how many categories is one-hot encoding typically appropriate?

- A) Unlimited categories
- B) Few categories (typically <15-20)
- C) Only 2 categories
- D) 100+ categories

ANSWER: B

Q9. What is the "dummy variable trap" in one-hot encoding?

- A) Using too few categories
- B) Multicollinearity when using all dummy variables
- C) Missing data problem
- D) Outlier detection issue

ANSWER: B

Q10. How do you avoid the dummy variable trap?

- A) Use all dummy variables
- B) Drop one dummy variable (first or last)
- C) Never use one-hot encoding
- D) Add more categories

ANSWER: B

Q11. What is target encoding (mean encoding)?

- A) Encoding the target as 0 and 1
- B) Replacing category with mean of target variable for that category
- C) Creating binary columns
- D) Using category counts

ANSWER: B

Q12. What is the main advantage of target encoding?

- A) Simplest to implement
- B) Captures relationship with target (high information)
- C) Never causes problems
- D) Works without target variable

ANSWER: B

Q13. What is the main risk of target encoding?

- A) Too slow
- B) Target leakage and overfitting
- C) Reduces accuracy
- D) Only works for regression

ANSWER: B

Q14. What is frequency encoding?

- A) Encoding based on alphabetical order
- B) Replacing category with its frequency/count
- C) Creating binary columns
- D) Using target mean

ANSWER: B

Q15. What is binary encoding used for?

- A) Only binary variables

B) High-cardinality features (converts to binary code)

C) Continuous variables

D) Ordinal variables only

ANSWER: B

LEVEL 2: FEATURE ENGINEERING & TRANSFORMATION (Questions 16-35)

Q16. What is feature engineering?

- A) Selecting existing features
- B) Creating new features from existing data using domain knowledge
- C) Removing outliers
- D) Scaling features

ANSWER: B

Q17. What can be extracted from date/time features?

- A) Nothing useful
- B) Year, month, day, day_of_week, hour, is_weekend, season
- C) Only the year
- D) Only the day

ANSWER: B

Q18. What features can be extracted from text data?

- A) Nothing quantifiable
- B) Length, word count, sentiment score, TF-IDF vectors
- C) Only word count
- D) Only length

ANSWER: B

Q19. What are ratio features in feature engineering?

- A) Random ratios
- B) Meaningful ratios like debt_to_income, price_per_sqft
- C) All possible ratios
- D) Only percentages

ANSWER: B

Q20. What are aggregation features?

- A) Random groupings
- B) Group-based statistics (mean, median, std by category)
- C) Sum of all features
- D) Only counts

ANSWER: B

Q21. What are interaction features?

- A) Random combinations
- B) Multiplying or combining features to capture relationships
- C) Adding all features
- D) Deleting features

ANSWER: B

Q22. What is polynomial feature transformation?

- A) Only linear features
- B) Creates polynomial combinations like x^2 , x^3 , $x_1 \cdot x_2 \cdot x_3$
- C) Removes features
- D) Sorts features

ANSWER: B

Q23. What is a disadvantage of polynomial features?

- A) Too simple

- B) Dramatically increases dimensionality
- C) Always improves accuracy
- D) Only works for regression

ANSWER: B

Q24. What is binning (discretization)?

- A) Removing features
- B) Converting continuous to categorical ranges
- C) Only for categorical data
- D) Scaling method

ANSWER: B

Q25. When is log transformation useful?

- A) For all data types
- B) Right-skewed distributions and reducing outlier impact
- C) Only for negative values
- D) Normal distributions

ANSWER: B

Q26. What does log transformation do to multiplicative relationships?

- A) Removes them
- B) Converts to additive relationships
- C) Makes them exponential
- D) No effect

ANSWER: B

Q27. What is Box-Cox transformation?

- A) A fixed transformation
- B) Automatically finds optimal power transformation
- C) Only for categorical data
- D) A feature selection method

ANSWER: B

Q28. What limitation does Box-Cox have?

- A) Too complex
- B) Only for strictly positive values
- C) Only for negative values
- D) Works for all data

ANSWER: B

Q29. What is Yeo-Johnson transformation?

- A) Same as Box-Cox
- B) Similar to Box-Cox but handles zero and negative values
- C) Only for positive values
- D) A scaling method

ANSWER: B

Q30. What does automated feature engineering using Featuretools do?

- A) Manual feature creation
- B) Automatic deep feature synthesis
- C) Only scales features
- D) Removes features

ANSWER: B

Q31. What role does domain knowledge play in feature engineering?

- A) Not important
- B) Essential for identifying meaningful interactions
- C) Only needed for simple problems
- D) Can be completely ignored

ANSWER: B

Q32. What is the BMI formula as an example of feature engineering?

- A) weight × height
- B) weight / height²
- C) weight + height
- D) weight - height

ANSWER: B

Q33. What are lagged features in time series?

- A) Future values
- B) Previous values used as predictors
- C) Random values
- D) Mean values

ANSWER: B

Q34. What are rolling window features?

- A) Fixed values
- B) Statistics computed over moving time windows
- C) Only for images
- D) Random samples

ANSWER: B

Q35. Why is feature engineering considered both art and science?

- A) It's completely random
- B) Requires creativity and domain knowledge plus technical skill
- C) Only uses algorithms
- D) No creativity needed

ANSWER: B

LEVEL 3: FEATURE SELECTION METHODS (Questions 36-55)

Q36. Why is feature selection important?

- A) Makes data look better
- B) Reduces dimensionality, overfitting, and improves performance
- C) Always worse for models
- D) Only for visualization

ANSWER: B

Q37. What are the three main categories of feature selection methods?

- A) Fast, medium, slow
- B) Filter, Wrapper, Embedded methods
- C) Simple, medium, complex
- D) Linear, nonlinear, mixed

ANSWER: B

Q38. What are filter methods in feature selection?

- A) Model-dependent methods
- B) Model-independent statistical tests
- C) Only for neural networks
- D) Require training models

ANSWER: B

Q39. What does variance threshold do?

- A) Adds variance
- B) Removes low-variance (near-constant) features
- C) Increases features
- D) Scales features

ANSWER: B

Q40. Why remove low-variance features?

- A) They're too important
- B) Near-constant values provide little information
- C) To increase complexity
- D) Required by law

ANSWER: B

Q41. What does correlation analysis identify?

- A) Independent features
- B) Highly correlated features (multicollinearity)
- C) Important features only
- D) Outliers

ANSWER: B

Q42. What is the chi-squared test used for in feature selection?

- A) Numerical features only
- B) Categorical features with categorical target
- C) Only continuous variables
- D) Model training

ANSWER: B

Q43. What is the ANOVA F-test used for?

- A) Categorical features only

- B) Numerical features with categorical target
- C) Only for regression
- D) Only for trees

ANSWER: B

Q44. What does mutual information measure?

- A) Feature size
- B) Dependency between features and target
- C) Only linear relationships
- D) Feature count

ANSWER: B

Q45. What is an advantage of filter methods?

- A) Considers interactions
- B) Fast, model-agnostic, no overfitting risk
- C) Always most accurate
- D) Requires many models

ANSWER: B

Q46. What is a disadvantage of filter methods?

- A) Too slow
- B) Ignores feature interactions (univariate)
- C) Always causes overfitting
- D) Requires target variable

ANSWER: B

Q47. What are wrapper methods?

- A) Model-independent methods
- B) Model-dependent methods that search feature subsets
- C) Only statistical tests
- D) Scaling methods

ANSWER: B

Q48. What is forward selection?

- A) Remove all features
- B) Start empty, iteratively add best feature
- C) Start with all features
- D) Random selection

ANSWER: B

Q49. What is backward elimination?

- A) Add features one by one
- B) Start with all, iteratively remove worst feature
- C) Random removal
- D) Keep all features

ANSWER: B

Q50. What is Recursive Feature Elimination (RFE)?

- A) Adds features recursively
- B) Recursively removes least important features
- C) Random selection
- D) One-time removal

ANSWER: B

Q51. What is an advantage of wrapper methods?

- A) Very fast
- B) Considers feature interactions
- C) Never overfits
- D) No model needed

ANSWER: B

Q52. What is a disadvantage of wrapper methods?

- A) Too simple
- B) Computationally expensive, overfitting risk
- C) Always fast
- D) Never accurate

ANSWER: B

Q53. What are embedded methods?

- A) Separate from training
- B) Built-in feature selection during model training
- C) Only filter methods
- D) Manual selection

ANSWER: B

Q54. How does Lasso (L1) perform feature selection?

- A) Keeps all features equally
- B) Drives some coefficients to exactly zero
- C) Only scales features
- D) Adds features

ANSWER: B

Q55. How do tree-based models provide feature importance?

- A) Random scores
- B) Based on split improvement and usage frequency
- C) All features equal
- D) External calculation

ANSWER: B

LEVEL 4: DIMENSIONALITY REDUCTION & TRAIN/TEST SPLITTING (Questions 56-70)

Q56. What is the difference between feature selection and feature extraction?

- A) No difference
- B) Selection keeps original features, extraction creates new ones
- C) Selection always better
- D) Extraction always better

ANSWER: B

Q57. What does PCA (Principal Component Analysis) do?

- A) Selects best features
- B) Creates new uncorrelated features capturing maximum variance
- C) Removes outliers
- D) Scales features

ANSWER: B

Q58. Are PCA components the same as original features?

- A) Yes, exactly the same
- B) No, linear combinations of original features
- C) Only sometimes
- D) Only for small datasets

ANSWER: B

Q59. What does LDA (Linear Discriminant Analysis) maximize?

- A) Variance only
- B) Class separability
- C) Feature count
- D) Training speed

ANSWER: B

Q60. What is t-SNE used for?

- A) Model training
- B) Non-linear dimensionality reduction for visualization
- C) Feature scaling
- D) Outlier detection

ANSWER: B

Q61. What is UMAP?

- A) Same as PCA
- B) Faster than t-SNE, preserves global structure
- C) Only for categorical data
- D) A classification algorithm

ANSWER: B

Q62. What are autoencoders?

- A) Manual encoders
- B) Neural network-based dimensionality reduction
- C) Only for images
- D) Statistical methods

ANSWER: B

Q63. Why split data into training and test sets?

- A) To increase data volume
- B) To evaluate model performance on unseen data
- C) To speed up training
- D) Required by Python

ANSWER: B

Q64. What is the typical train/test split ratio?

- A) 50/50
- B) 80/20 or 70/30
- C) 95/5
- D) 99/1

ANSWER: B

Q65. What is stratified sampling?

- A) Random sampling only
- B) Maintains class distribution in all splits
- C) Only for large datasets
- D) Removes rare classes

ANSWER: B

Q66. When is stratified sampling essential?

- A) Always, for all problems
- B) For imbalanced classification datasets
- C) Never needed
- D) Only for regression

ANSWER: B

Q67. What is a train/validation/test split?

- A) Two-way split
- B) Three-way split (e.g., 60/20/20) with validation for tuning
- C) No splitting needed
- D) Only train and test

ANSWER: B

Q68. What is the validation set used for?

- A) Final evaluation
- B) Hyperparameter tuning
- C) Initial training
- D) Data cleaning

ANSWER: B

Q69. What is the test set used for?

- A) Hyperparameter tuning
- B) Final unbiased evaluation (use only once)
- C) Training the model
- D) Feature selection

ANSWER: B

Q70. What is time series splitting?

- A) Random splitting
- B) Sequential validation respecting temporal order
- C) Any split method
- D) Only for images

ANSWER: B

LEVEL 5: CROSS-VALIDATION TECHNIQUES (Questions 71-85)

Q71. What is cross-validation?

- A) Single train/test split
- B) Multiple train/test splits for robust performance estimates
- C) Only for final testing
- D) Data cleaning method

ANSWER: B

Q72. What is K-Fold cross-validation?

- A) One split
- B) Split data into K equal folds, train on K-1, test on 1, repeat K times
- C) Random splits
- D) No splitting

ANSWER: B

Q73. What are typical values for K in K-Fold CV?

- A) $K = 2$
- B) $K = 5$ or $K = 10$
- C) $K = 100$
- D) $K = 1$

ANSWER: B

Q74. How is the final cross-validation score calculated?

- A) Use only best fold
- B) Average performance across all K folds
- C) Use only worst fold
- D) Random selection

ANSWER: B

Q75. What is stratified K-Fold?

- A) Same as regular K-Fold
- B) Maintains class distribution in each fold
- C) Only for regression
- D) Random folding

ANSWER: B

Q76. When is stratified K-Fold essential?

- A) Never needed
- B) For imbalanced classification problems
- C) Only for large datasets
- D) All regression problems

ANSWER: B

Q77. What is Leave-One-Out Cross-Validation (LOOCV)?

- A) Leave out 50% of data
- B) $K = n$ (each sample is one fold)
- C) Random leaving out
- D) No leaving out

ANSWER: B

Q78. When is LOOCV appropriate?

- A) Large datasets (10,000+ samples)

- B) Very small datasets (<100 samples)
- C) Always the best choice
- D) Never appropriate

ANSWER: B

Q79. What is repeated K-Fold CV?

- A) Single K-Fold run
- B) Repeat K-Fold multiple times with different random splits
- C) No repetition
- D) Only two repetitions

ANSWER: B

Q80. What is time series cross-validation?

- A) Random CV
- B) Forward-chaining with progressively increasing training window
- C) Standard K-Fold
- D) No CV for time series

ANSWER: B

Q81. Why can't you use standard K-Fold for time series?

- A) Too slow
- B) Would use future data in training (look-ahead bias)
- C) Works perfectly fine
- D) Only technical limitation

ANSWER: B

Q82. What is Group K-Fold?

- A) Random grouping
- B) Ensures groups (e.g., patients) not split across folds
- C) Only for classification
- D) Same as stratified

ANSWER: B

Q83. When is Group K-Fold needed?

- A) Never
- B) When multiple samples from same source (prevent leakage)
- C) Always for all problems
- D) Only for regression

ANSWER: B

Q84. What is nested cross-validation?

- A) Single CV loop
- B) Outer loop for evaluation, inner loop for hyperparameter tuning
- C) No nesting
- D) Only for neural networks

ANSWER: B

Q85. Why is nested CV considered the gold standard?

- A) Fastest method
- B) Provides unbiased performance estimate
- C) Easiest to implement
- D) Requires no computation

ANSWER: B

LEVEL 6: CLASSIFICATION METRICS - CONFUSION MATRIX (Questions 86-100)

Q86. What does a confusion matrix show?

- A) Feature correlations
- B) Actual vs predicted classifications
- C) Training time
- D) Feature importance

ANSWER: B

Q87. What are True Positives (TP)?

- A) Predicted positive, actually negative
- B) Predicted positive, actually positive (correct)
- C) Predicted negative, actually positive
- D) Predicted negative, actually negative

ANSWER: B

Q88. What are False Positives (FP)?

- A) Predicted positive, actually positive
- B) Predicted positive, actually negative (Type I Error)
- C) Predicted negative, actually positive
- D) Predicted negative, actually negative

ANSWER: B

Q89. What are False Negatives (FN)?

- A) Predicted positive, actually negative
- B) Predicted negative, actually negative
- C) Predicted negative, actually positive (Type II Error)
- D) Predicted positive, actually positive

ANSWER: C

Q90. What are True Negatives (TN)?

- A) Predicted positive, actually negative
- B) Predicted negative, actually positive
- C) Predicted positive, actually positive
- D) Predicted negative, actually negative (correct)

ANSWER: D

Q91. What is a Type I Error?

- A) False Negative
- B) False Positive
- C) True Positive
- D) True Negative

ANSWER: B

Q92. What is a Type II Error?

- A) False Positive
- B) False Negative
- C) True Positive
- D) True Negative

ANSWER: B

Q93. What is accuracy in classification?

- A) TP / (TP + FP)

- B) $(TP + TN) / (TP + TN + FP + FN)$
- C) $TP / (TP + FN)$
- D) $TN / (TN + FP)$

ANSWER: B

Q94. Why is accuracy misleading for imbalanced datasets?

- A) It's always accurate
- B) Can achieve high accuracy by predicting majority class
- C) Only works for balanced data technically
- D) Too complex to calculate

ANSWER: B

Q95. If 95% of emails are not spam, what accuracy can you get by always predicting "not spam"?

- A) 50%
- B) 95% (without learning anything useful)
- C) 5%
- D) 0%

ANSWER: B

Q96. What is precision in classification?

- A) $TP / (TP + FN)$
- B) $TP / (TP + FP)$ - Of predicted positives, how many actually positive
- C) $TN / (TN + FP)$
- D) $(TP + TN) / \text{Total}$

ANSWER: B

Q97. What question does precision answer?

- A) Of actual positives, how many did we catch?
- B) When model says positive, how often is it right?
- C) What's the overall accuracy?
- D) How many negatives are there?

ANSWER: B

Q98. What is recall (sensitivity)?

- A) $TP / (TP + FP)$
- B) $TP / (TP + FN)$ - Of actual positives, how many correctly identified
- C) $TN / (TN + FP)$
- D) $(TP + TN) / \text{Total}$

ANSWER: B

Q99. What question does recall answer?

- A) When model says positive, how often right?
- B) Of all actual positive cases, how many did model catch?
- C) What's the overall accuracy?
- D) How many false positives?

ANSWER: B

Q100. Which metric focuses on minimizing false positives?

- A) Recall
- B) Precision
- C) Accuracy
- D) Specificity

ANSWER: B

END OF FIRST 100 QUESTIONS (Pages 12+)

NOTE: This is Part 1 of the questions covering pages 12+.

Additional 100 questions would be needed to cover:

SCORE INTERPRETATION:

90-100: Excellent understanding of encoding, feature engineering, and evaluation

80-89: Very good grasp with minor gaps

70-79: Good foundation, review feature selection methods

60-69: Fair understanding, focus on CV and metrics

Below 60: Review encoding methods and feature engineering basics

STUDY FOCUS: